

Data Preprocessing

first i set option to display all columns

for less memory usage and faster training i converted dtype from float64 to float32 and int64 to int8

checking for any missing value

checking for any duplicated row

deleting single-valued column bcz it will not contribute anything to model prediction

checking min and max value of each column

separating output labels from features

Identifying and Handling Outliers

Outliers are very big problem in this dataset
Some of columns contain values in the order from $e+0$ to $e+9$ which will definitely affect model performance

first i tried IQR method but it identifies most of rows as outlier so this is not good method here then

i took log of columns which contain very extreme values
taking $\log(1+x)$ so that very small values does not get converted to -infinity

Robust Scaling

it handles outliers by scaling based on median and IQR

Visualizations

Class distributions

Now biggest problem is that data is highly imbalanced, about only 3% cases are bankrupt so we will not prefer accuracy much here

![Class Distributions Plot][images/class_distributions.png]

Correlation Heatmap

hello this is correlation heatmap

splitting dataset into training and testing set

SMOTE (Synthetic minority oversampling technique)

In this imbalanced dataset we have two options either we can use `class_weight="balanced"` which works in most of the models or we can balance dataset by oversampling minority class

At first glance i thought that all columns of data are scaled and also didn't thought about outliers and used `class_weight="balanced"` (No SMOTE) and trained various models

I will share their results in very last for just comparison how things get changed

Now we train various models on this preprocessed and balanced dataset

Logistic Regression

(default) `solver=lbfgs, penalty=l2`

	precision	recall	f1-score	support
0	0.99	0.88	0.94	1324
1	0.18	0.82	0.29	40

Confusion Matrix: $\begin{bmatrix} 1171 & 153 \\ 7 & 33 \end{bmatrix}$

ROC AUC Score: 0.93

PR AUC score: 0.30

Accuracy_score: 0.88

best_f1: 0.42

best_thresh: 0.81

`solver="liblinear", penalty="l2"`

	precision	recall	f1-score	support
0	0.99	0.88	0.94	1324
1	0.17	0.80	0.28	40

Confusion Matrix: $\begin{bmatrix} 1170 & 154 \\ 8 & 32 \end{bmatrix}$

ROC AUC Score: 0.93

PR AUC score: 0.30

Accuracy_score: 0.88

best_f1: 0.41

best_thresh: 0.80

`solver="saga", penalty="l1"`

	precision	recall	f1-score	support
0	0.99	0.88	0.94	1324
1	0.17	0.80	0.28	40

Confusion Matrix: $\begin{bmatrix} 1171 & 153 \\ 8 & 32 \end{bmatrix}$

ROC AUC Score: 0.93

PR AUC score: 0.30

Accuracy_score: 0.88

best_f1: 0.42

best_thresh: 0.81

`solver="liblinear",penalty="l1"`

	precision	recall	f1-score	support
0	0.99	0.89	0.94	1324
1	0.18	0.78	0.29	40

Confusion Matrix: $\begin{bmatrix} 1179 & 145 \\ 9 & 31 \end{bmatrix}$

ROC AUC Score: 0.93

PR AUC score: 0.31

Accuracy_score: 0.89

best_f1: 0.38

best_thresh: 0.86

`best_f1` for logistic regression is 0.42

`best model= (default) solver=lbfgs, penalty=l2`

Regularization

`C:[0.01,0.1,1,5,10]`

best_f1: 0.30, best_thresh: 0.64

best_f1: 0.37, best_thresh: 0.72

best_f1: 0.41, best_thresh: 0.80

best_f1: 0.39, best_thresh: 0.78

best_f1: 0.40, best_thresh: 0.77

`C: [0.5,1,1.5,2,2.5,3]`

best_f1: 0.41, best_thresh: 0.77

best_f1: 0.41, best_thresh: 0.80

best_f1: 0.42, best_thresh: 0.81

```
best_f1: 0.41, best_thresh: 0.81
best_f1: 0.41, best_thresh: 0.82
best_f1: 0.40, best_thresh: 0.82
```

best regularization strength= 1.5

	precision	recall	f1-score	support
0	0.99	0.89	0.94	1324
1	0.17	0.78	0.28	40

```
Confusion Matrix: [[1176 148]
                   [  9  31]]
```

ROC AUC Score: 0.93

PR AUC score: 0.30

Accuracy_score: 0.88

best_f1: 0.42

best_thresh: 0.81

but our precision and recall are lower than default one, thus we will stick with default solver=lbfgs, penalty=l2, C=1

Support Vector Machines

Training on imbalanced dataset

```
svm_clf=SVC(
    kernel="rbf",
    probability=True,
    class_weight="balanced",
    C=1.0,
    gamma="scale"
)
```

	precision	recall	f1-score	support
0	0.98	0.66	0.79	1324
1	0.05	0.57	0.09	40

```
Confusion Matrix: [[872 452]
                   [ 17  23]]
```

ROC AUC Score: 0.73

PR AUC score: 0.10

Accuracy_score: 0.66

best_f1: 0.17

best_thresh: 0.06

Here SVM tries to draw a boundary that maximizes margin but with very few minority class samples it heavily favours majority class

On balanced dataset

	precision	recall	f1-score	support
0	0.98	0.76	0.86	1324
1	0.06	0.50	0.11	40

Confusion Matrix: $\begin{bmatrix} 1006 & 318 \\ 20 & 20 \end{bmatrix}$

ROC AUC Score: 0.71

PR AUC score: 0.06

Accuracy_score: 0.75

best_f1: 0.13

best_thresh: 0.83

Here balanced dataset performed better but after threshold optimizing imbalanced dataset has better score, but we would focus on balanced dataset bcz it gives a more balanced decision boundary and also helpful for recall

GridSearchCV

```
params={
    "C": [0.1, 0.5, 1, 2],
    "kernel": ["rbf"],
    "gamma": ["scale"]
}
```

Decision Tree

(default) max_depth=None

	precision	recall	f1-score	support
0	0.98	0.95	0.97	1324
1	0.23	0.45	0.31	40

Confusion Matrix: $\begin{bmatrix} 1264 & 60 \\ 22 & 18 \end{bmatrix}$

ROC AUC Score: 0.70

PR AUC score: 0.12

Accuracy_score: 0.94

best_f1: 0.31

best_thresh: 0.01

```

max_depth=3, best_f1: 0.35
            best_thresh: 0.82
max_depth=5, best_f1: 0.36
            best_thresh: 0.94
max_depth=6, best_f1: 0.38
            best_thresh: 0.94
max_depth=7, best_f1: 0.36
            best_thresh: 0.96
max_depth=10, best_f1: 0.33
            best_thresh: 0.51

```

best_f1 for decision tree is 0.38

best model= (max_depth=6)

RandomForestClassifier

(max_depth=6)

```

n_estimators=100, best_f1: 0.48
                best_thresh: 0.84

```

```

n_estimators=150, best_f1: 0.49
                best_thresh: 0.84

```

precision	recall	f1-score	support	
0	0.99	0.92	0.95	1324
1	0.22	0.78	0.34	40

```

Confusion Matrix: [[1213  111]
                   [   9   31]]

```

ROC AUC Score: 0.94

PR AUC score: 0.42

Accuracy_score: 0.91

best_f1: 0.49

best_thresh: 0.85

```

n_estimators=200, best_f1: 0.48
                best_thresh: 0.86

```

(max_depth=8, n_estimators=150)

	precision	recall	f1-score	support
0	0.99	0.94	0.96	1324
1	0.26	0.75	0.38	40

Confusion Matrix: $\begin{bmatrix} 1238 & 86 \\ 10 & 30 \end{bmatrix}$

ROC AUC Score: 0.94

PR AUC score: 0.42

Accuracy_score: 0.93

best_f1: 0.48

best_thresh: 0.81

best_f1 for RandomForestClassifier is 0.48

best model= (max_depth=6, n_estimators=150)

GridSearchCV

MY MISTAKE HERE- Since SMOTE is applied before and cross-validation is used in GridSearch, we would train on synthetic data which is easier and does not represent real world performance thus

We perform GridSearchCV on dataset with (No SMOTE)

```
params={
    "max_depth": [6, 8, 10],
    "max_features": ["sqrt", "log2"],
    "n_estimators": [100, 150, 200],
    "bootstrap": [True],
}
```

```
{'bootstrap': True, 'max_depth': 8, 'max_features': 'log2', 'n_estimators': 150}
grid.best_score_ = 0.45
```

	precision	recall	f1-score	support
0	0.99	0.96	0.98	1324
1	0.33	0.60	0.43	40

Confusion Matrix: $\begin{bmatrix} 1276 & 48 \\ 16 & 24 \end{bmatrix}$

ROC AUC Score: 0.94

PR AUC score: 0.45

Accuracy_score: 0.95

best_f1: 0.53

best_thresh: 0.62

Again with more hyperparameters

```
params={
    "max_depth": [6, 8, 10],
```

```

    "max_features":["sqrt","log2"],
    "n_estimators":[150,200],
    "bootstrap":[True],
    "min_samples_split":[10,20,40],
    "min_samples_leaf":[5,10,15]
}

{'bootstrap': True, 'max_depth': 8, 'max_features': 'sqrt', 'min_samples_leaf': 10, 'min_samples_split': 20, 'n_estimators': 200}
grid.best_score_: 0.479

```

	precision	recall	f1-score	support
0	0.99	0.95	0.97	1324
1	0.29	0.68	0.41	40

```

Confusion Matrix: [[1259  65]
                   [  13  27]]
ROC AUC Score: 0.94
PR AUC score: 0.43
Accuracy_score: 0.94
best_f1: 0.50
best_thresh: 0.70

```

this time results were lower but we have used min_samples_split and min_samples_leaf

Again with slightly different parameter space

```

params={
    "max_depth":[8],
    "max_features":["sqrt","log2"],
    "n_estimators":[150,200],
    "bootstrap":[True],
    "min_samples_split":[15,20,25,30],
    "min_samples_leaf":[7,10,13]
}

{'bootstrap': True, 'max_depth': 8, 'max_features': 'sqrt', 'min_samples_leaf': 7, 'min_samples_split': 15, 'n_estimators': 150}
grid.best_score_ = 0.47

```

	precision	recall	f1-score	support
0	0.99	0.96	0.97	1324
1	0.32	0.68	0.44	40

```

Confusion Matrix: [[1267  57]
                   [  13  27]]
ROC AUC Score: 0.94
PR AUC score: 0.49
Accuracy_score: 0.95
best_f1: 0.52
best_thresh: 0.66

```


Again tweaking some values

```
params={
    "max_depth":[8],
    "max_features":["sqrt","log2"],
    "n_estimators":[150,200],
    "bootstrap":[True],
    "min_samples_split":[12,15,18],
    "min_samples_leaf":[3,5,7,10]
}
```

XGBoostClassifier

```
n_estimators=500,
learning_rate=0.05,
max_depth=5,
scale_pos_weight=scale_pos_weight,
eval_metric="aucpr",
early_stopping_rounds=30
```

	precision	recall	f1-score	support
0	0.99	0.94	0.97	1324
1	0.27	0.70	0.39	40

```
Confusion Matrix: [[1249   75]
                   [  12   28]]
```

```
ROC AUC Score: 0.91
```

```
PR AUC score: 0.39
```

```
Accuracy_score: 0.94
```

```
best_f1: 0.45
```

```
best_thresh: 0.81
```

GridSearchCV

```
params={
    'n_estimators':[300],
    'max_depth':[3,5,7],
    'gamma':[0,0.1,0.3],
    'learning_rate':[0.01,0.05,0.1],
    'min_child_weight':[1,5,10]
}
```

```
{'gamma': 0, 'learning_rate': 0.1, 'max_depth': 5, 'min_child_weight': 5,
'n_estimators': 300}
grid.best_score_ = 0.51
```

	precision	recall	f1-score	support
0	0.99	0.98	0.98	1324
1	0.45	0.53	0.48	40

Confusion Matrix: $\begin{bmatrix} 1298 & 26 \\ 19 & 21 \end{bmatrix}$

ROC AUC Score: 0.93
 PR AUC score: 0.44
 Accuracy_score: 0.97
 best_f1: 0.54
 best_thresh: 0.40

Again GridSearchCV

```
params={
    'n_estimators':[300],
    'max_depth':[4,5,6],
    'gamma':[0,0.2],
    'learning_rate':[0.075,0.1,0.125],
    'subsample':[0.8,1],
    'colsample_bytree':[0.8,1],
    'min_child_weight':[5,7]
}
```

```
{'colsample_bytree': 0.8, 'learning_rate': 0.125, 'max_depth': 6,
'min_child_weight': 7, 'n_estimators': 300, 'subsample': 1}
grid.best_score_ = 0.525
```

	precision	recall	f1-score	support
0	0.99	0.98	0.98	1324
1	0.46	0.55	0.50	40

Confusion Matrix: $\begin{bmatrix} 1298 & 26 \\ 18 & 22 \end{bmatrix}$

ROC AUC Score: 0.93
 PR AUC score: 0.41
 Accuracy_score: 0.97
 best_f1: 0.53
 best_thresh: 0.40

**this time best_f1 after threshold optimization and PR AUC score are lower than previous one 0.54 so
 We will choose best model from previous grid search**